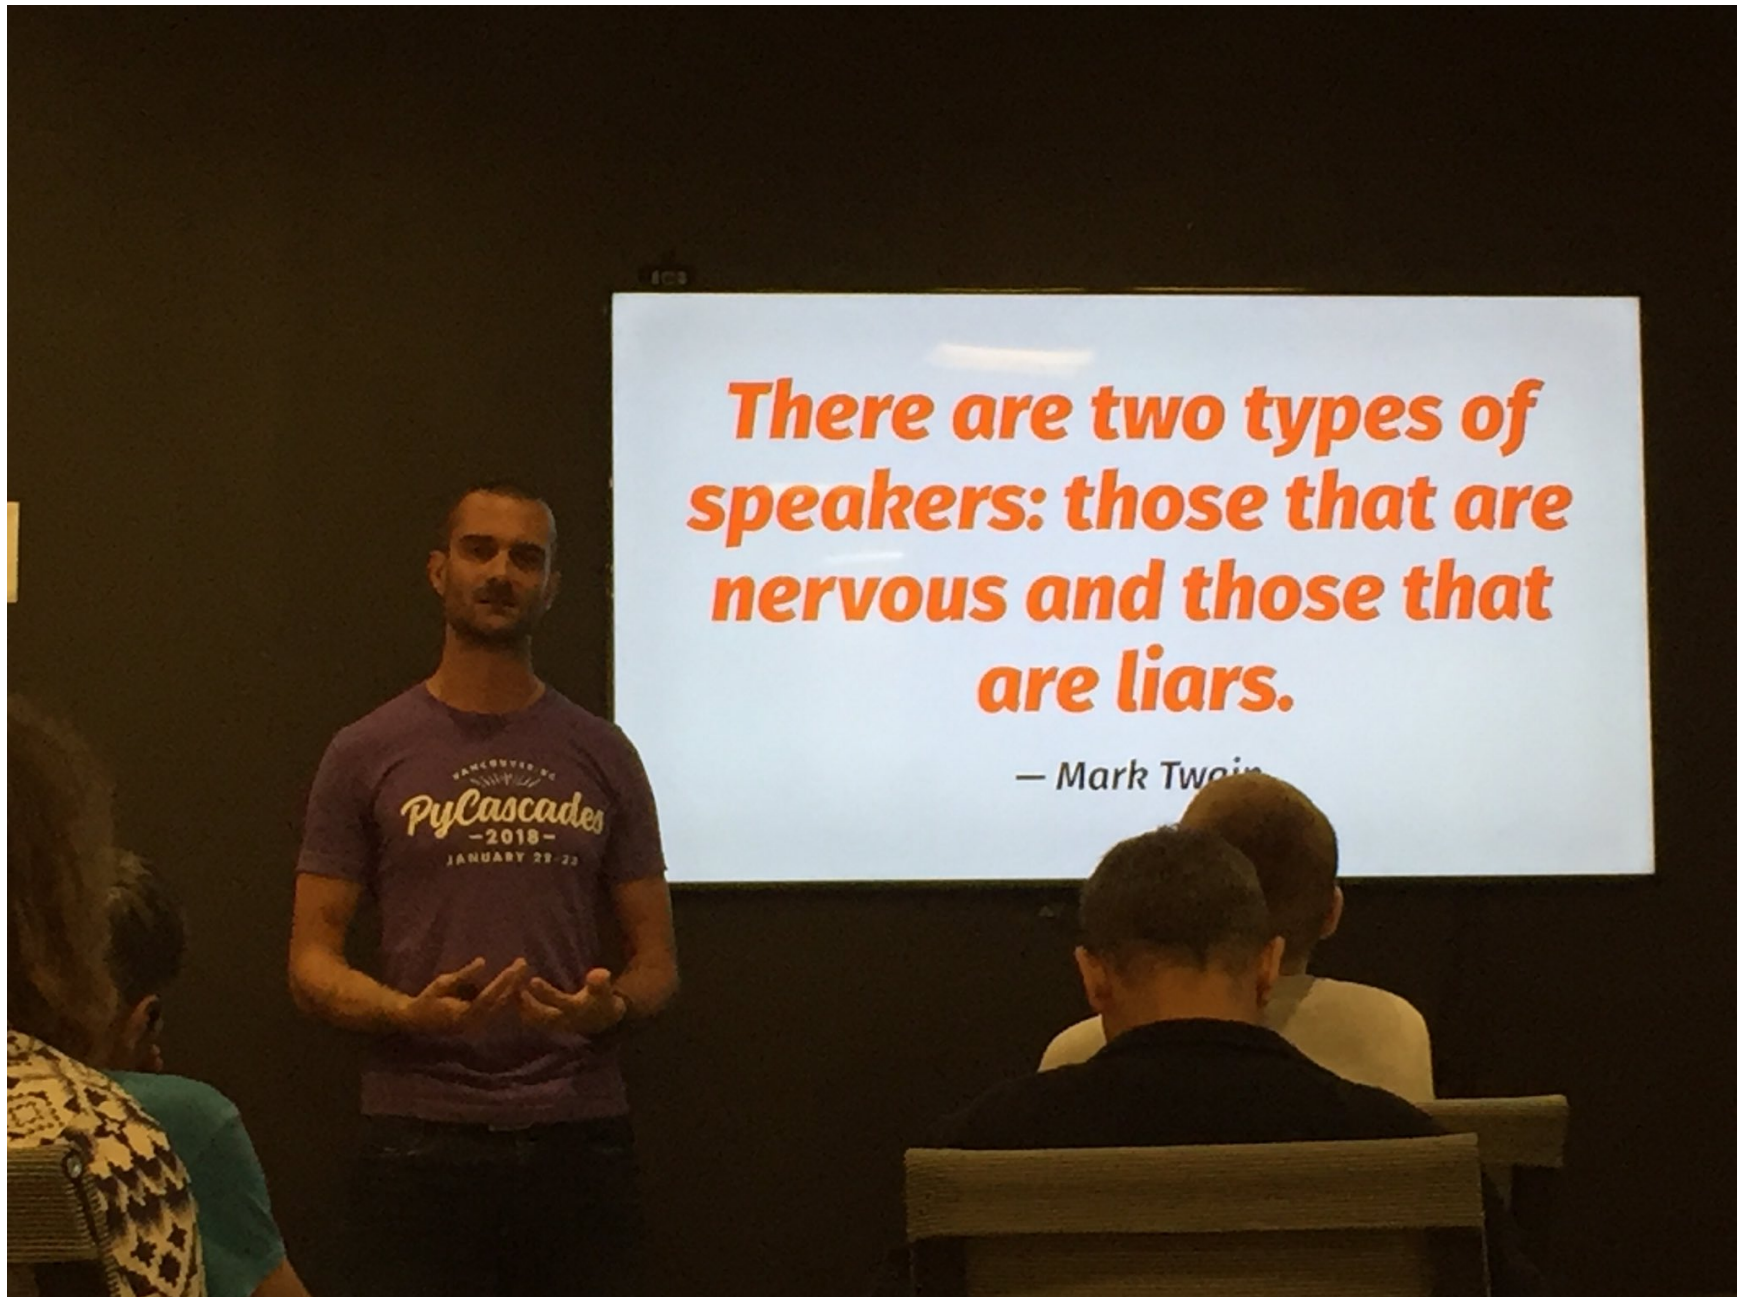


THE MODERN PROMETHEUS

PIOTR SZOTKOWSKI



DAN BADER'S PHOTO OF SEBASTIAN VETTER'S TALK



JERZY STRZELECKI, SCHLOSS FRANKENSTEIN

FRANKENSTEIN ;
OR,
THE MODERN PROMETHEUS.



IN THREE VOLUMES.



(DARMSTADT, HESSEN)

FRANKENSTEIN PROGRAMMING

AS IN

~ SEAMLESSLY
COMBINING MORE
THAN ONE LANGUAGE




John T. McInnes

@lyebot

Common misconception; "Big Ben" is actually the name of the clock's creator. Its correct name is "Big Ben's Monster"



BENCHMARKS



Synthetic benchmarks tell us sweet FA about real world performance of code, **architecture** being a much more significant consideration than the proportion of raw MIPS a given language will deliver on a given platform.

The average netbook could happily run all of Teller's fusion bomb models along with the full telemetry analysis of all the Apollo missions in the pauses between loading XKCD comics and binning junk mail without the user being any the wiser.

— Eleanor McHugh

```
require 'active_record'
require 'benchmark/ips'
```

```
ActiveRecord::Base.establish_connection adapter: 'postgresql'
```

```
class CreateWhatevers < ActiveRecord::Migration[5.2]
  def change
    create_table(:whatevers) { |table| table.text :text }
  end
end
```

```
CreateWhatevers.migrate :up
at_exit { CreateWhatevers.migrate :down }
```

```
Whatever = Class.new(ActiveRecord::Base)
```

```
Benchmark.ips do |bench|
  bench.report('exception miss') { raise '💣' if false }
  bench.report('exception hit') { raise '💣' rescue nil }
  bench.report('SELECT') { Whatever.first }
  bench.report('INSERT') { Whatever.create(text: 'meh') }
  bench.compare!
end
```


Warming up -----

exception miss	445.207k	i/100ms
exception hit	107.302k	i/100ms
SELECT	546.000	i/100ms
INSERT	23.000	i/100ms

Calculating -----

exception miss	13.660M	(± 3.6%)	i/s -	68.562M	in	5.026944s
exception hit	1.375M	(± 2.3%)	i/s -	6.975M	in	5.074398s
SELECT	5.108k	(± 4.5%)	i/s -	25.662k	in	5.034214s
INSERT	241.202	(± 9.1%)	i/s -	1.219k	in	5.090564s

Comparison:

exception miss:	13660293.4	i/s	
exception hit:	1375266.0	i/s - 9.93x	slower
SELECT:	5108.2	i/s - 2674.19x	slower
INSERT:	241.2	i/s - 56634.16x	slower



Isaac Sloan

@elorest

968,824.35 req/sec: 32 cores at 2.7Ghz
#AmberFramework on @CrystalLanguage. I
love @rubyonrails but it and @elixirphoenix
seem like turtles.

```
ubuntu@ip-172-31-0-70:~/bench> wrk -d 60 -t 20 -c 1015 http://localhost:3000
Running 1m test @ http://localhost:3000
20 threads and 1015 connections
Thread Stats   Avg      Stdev     Max   +/-  Stdev
Latency       1.86ms    2.88ms   56.61ms  87.54%
Req/Sec      48.73k    6.01k    88.40k   68.28%
58225168 requests in 1.00m, 4.01GB read
Requests/sec: 968824.35
Transfer/sec:   68.37MB
```



Tweet Quotes

@QuoteTwitts

If we have data, let's look at data. If all we have are opinions, let's go with mine. 💋 Jim Barksdale #Quotes

1. ALWAYS BENCHMARK

2. ADJUST ALGORITHM

3. IMPLEMENTATIONS?

(BUT IF YOU DID
AND IT'S *STILL* TOO SLOW
IT MIGHT BE WORTHWHILE
TO LOOK INTO OTHER LANGUAGES...)



Jamie Gaskins

@jamie_gaskins

Pivotal is the company that created the story-tracking software. The actual name of the software is Pivotal's Monster.

BINARY REPRESENTATION

LET'S HAVE A SET OF ELEMENTS (E.G., INTEGERS)

LET'S REPRESENT IT BY 'SET' BITS
(ONES IN ITS BINARY REPRESENTATION)

$$\{0, 2, 3, 5\} \rightarrow 0b101101$$

SET OPERATIONS ARE NOW BINARY OPERATIONS

$$\{0, 2, 3, 5\} \cap \{2, 3, 4\} \rightarrow 0b101101 \ \& \ 0b011100$$

$$0b101101 \ \& \ 0b011100 \rightarrow 0b001100 \rightarrow \{2, 3\}$$

POPULATION COUNT

LET'S SAY WE NEED THE SIZE OF OUR SET

NUMBER OF ONES IN BINARY REPRESENTATION

$$|\{0,2,3,5\}| \rightarrow \text{popcount}(0b101101) \rightarrow 4$$

$$|\{2,3,4\}| \rightarrow \text{popcount}(0b011100) \rightarrow 3$$

$$|\{2,3\}| \rightarrow \text{popcount}(0b001100) \rightarrow 2$$

LET'S FREEDOM-PATCH Integer

```
class Integer

  def popcount_to_s
    to_s(2).count('1')
  end

  def popcount_cont_shift
    count = 0
    number = self
    until number.zero?
      count += number & 1
      number >>= 1
    end
    count
  end
end
```

```
class Integer
```

```
  def popcount_bit_elim
```

```
    count = 0
```

```
    number = self
```

```
    until number.zero?
```

```
      number &= number - 1
```

```
      count += 1
```

```
    end
```

```
    count
```

```
  end
```

```
  def popcount_prog_shift
```

```
    number = self
```

```
    number -= (number >> 1) & 0x5555555555555555
```

```
    number = (number & 0x3333333333333333) + ((number >> 2) & 0x3333333333333333)
```

```
    number = (number + (number >> 4)) & 0x0f0f0f0f0f0f0f0f
```

```
    number = (number + (number >> 8)) & 0x00ff00ff00ff00ff
```

```
    number = (number + (number >> 16)) & 0x0000ffff0000ffff
```

```
    number += number >> 32
```

```
    (number + (number >> 64)) & 0xff
```

```
  end
```

```
end
```

RUST VIA HELIX? SURE!

```
#[macro_use]
extern crate helix;

ruby! {
  class Popcount {
    def count(int: u64) -> u32 {
      int.count_ones()
    }
  }
}
```

C? SSE4.2? WHY NOT!

```
class Integer
```

```
    inline do |builder|
      builder.c 'int popcount_bit_elim_c() {
        long number = NUM2LONG(self);
        int count;
        for (count = 0; number; count++) number &= number - 1;
        return count;
      }'
    end
```

```
    inline do |builder|
      builder.c 'int popcount_builtin() {
        return __builtin_popcountl(NUM2LONG(self));
      }'
    end
```

```
end
```

ALGORITHMS? APPROACH!

```
class Integer
```

```
  POPCOUNT_CACHE = (0x0000..0xffff).map { |number| number.to_s(2).count('1') }
```

```
  def popcount_cached
```

```
    POPCOUNT_CACHE[self & 0xffff] +  
    POPCOUNT_CACHE[self >> 16 & 0xffff] +  
    POPCOUNT_CACHE[self >> 32 & 0xffff] +  
    POPCOUNT_CACHE[self >> 48]
```

```
  end
```

```
end
```


LET'S BENCH! THAT INTROSPECTION!

```
require 'benchmark/ips'

numbers = Array.new(1_000) { rand(2**62) }
methods = Integer.instance_methods.grep(/^popcount_/)

raise 'oops' unless methods.map { |meth| numbers.map(&meth) }.uniq.size == 1

Benchmark.ips do |bench|
  methods.each do |meth|
    bench.report(meth[9..-1]) { numbers.map(&meth) }
  end
  bench.compare!
end
```

Ruby 2.5:

builtin:	29915.7 i/s		
bit_elim_c:	17011.6 i/s	- 1.76x	slower
cached:	4874.4 i/s	- 6.14x	slower
rust:	3133.3 i/s	- 9.55x	slower
prog_shift:	2878.5 i/s	- 10.39x	slower
to_s:	1131.3 i/s	- 26.44x	slower
bit_elim:	714.5 i/s	- 41.87x	slower
cont_shift:	317.3 i/s	- 94.27x	slower

JRuby 9.1:

cached:	5700.0 i/s		
prog_shift:	4894.2 i/s	- 1.16x	slower
bit_elim:	1337.3 i/s	- 4.26x	slower
to_s:	870.2 i/s	- 6.55x	slower
cont_shift:	716.2 i/s	- 7.96x	slower

JRuby 9.1 + 30 s warm-up + invoke dynamic:

prog_shift:	6749.9 i/s		
cached:	6566.6 i/s	- same-ish:	difference falls within error
bit_elim:	2587.5 i/s	- 2.61x	slower
cont_shift:	1644.5 i/s	- 4.10x	slower
to_s:	855.6 i/s	- 7.89x	slower

```
require "benchmark"
```

```
struct Int
```

```
  POPCOUNT_CACHE = (0x0000..0xffff).map { |number| number.to_s(2).count('1') }
```

```
  def popcount_bit_elim      # ...
```

```
  def popcount_cached       # ...
```

```
  def popcount_cont_shift   # ...
```

```
  def popcount_prog_shift   # ...
```

```
  def popcount_to_s         # ...
```

```
end
```

```
numbers = Array.new(1_000) { rand(2u64**62) }
```

```
Benchmark.ips do |bench|
```

```
  bench.report("popcount") { numbers.map(&.popcount) }
```

```
  bench.report("prog_shift") { numbers.map(&.popcount_prog_shift) }
```

```
  bench.report("cached") { numbers.map(&.popcount_cached) }
```

```
  bench.report("bit_elim") { numbers.map(&.popcount_bit_elim) }
```

```
  bench.report("cont_shift") { numbers.map(&.popcount_cont_shift) }
```

```
  bench.report("to_s") { numbers.map(&.popcount_to_s) }
```

```
end
```



RUBY TIMES INCLUDE RUBY-SIDE CALL OVERHEAD



Ruby 2.5 (+ Rust, C, SSE4.2...):

builtin:	29915.7 i/s		
bit_elim_c:	17011.6 i/s	- 1.76x	slower
cached:	4874.4 i/s	- 6.14x	slower
rust:	3133.3 i/s	- 9.55x	slower
prog_shift:	2878.5 i/s	- 10.39x	slower
to_s:	1131.3 i/s	- 26.44x	slower
bit_elim:	714.5 i/s	- 41.87x	slower
cont_shift:	317.3 i/s	- 94.27x	slower

Crystal 0.26:

popcount	267.47k	(3.74µs)	(± 1.13%)	8048 B/op		fastest
prog_shift	211.22k	(4.73µs)	(± 1.26%)	8048 B/op	1.27x	slower
cached	149.92k	(6.67µs)	(± 3.60%)	4048 B/op	1.78x	slower
bit_elim	33.96k	(29.45µs)	(± 1.12%)	4048 B/op	7.88x	slower
cont_shift	26.54k	(37.68µs)	(± 0.90%)	4048 B/op	10.08x	slower
to_s	8.14k	(122.78µs)	(± 3.15%)	84159 B/op	32.84x	slower



Lord Pinky

@HiddenPinky

"Cookie" is the name of the scientist who created him. It's Cookie's Monster, not "Cookie Monster". Read a book sometime.

ESCAPE ENTITIES

MAKE TEXT 'HTML SAFE'

< → <

> → >

& → &

' → '

" → "


```
class String
```

```
  ENTITIES = { '<' => '&lt;', '>' => '&gt;', '&' => '&amp;',  
               "'" => '&apos;', '"' => '&quot;' }
```

```
  def escape_each_char
```

```
    ''.tap do |result|
```

```
      each_char { |char| result << ENTITIES.fetch(char, char) }
```

```
    end
```

```
end
```

```
  def escape_gsub
```

```
    gsub(/[<>&'"]/ , ENTITIES)
```

```
end
```

```
  def escape CGI
```

```
    CGI.escape_html(self)
```

```
end
```

```
end
```

ZERO-COST ABSTRACTIONS...



```
#[macro_use]
extern crate helix;

ruby! {
  class HTML {
    def map_escape(input: String) -> String {
      input.chars().map(|chr| {
        match chr {
          '<' => String::from("&lt;"),
          '>' => String::from("&gt;"),
          '&' => String::from("&amp;"),
          '\'' => String::from("&apos;"),
          '"' => String::from("&quot;"),
          _   => chr.to_string(),
        }
      }).collect()
    }
  }
}
```

```

#[macro_use]
extern crate helix;

ruby! {
    class HTML {
        def push_escape(input: String) -> String {
            let mut result = String::with_capacity(2 * input.len());
            for chr in input.chars() {
                match chr {
                    '<' => result.push_str("&lt;"),
                    '>' => result.push_str("&gt;"),
                    '&' => result.push_str("&amp;"),
                    '\'' => result.push_str("&apos;"),
                    '"' => result.push_str("&quot;"),
                    _    => result.push(chr),
                }
            }
            result
        }
    }
}

```

Ruby 2.5:

rust_push:	286450.1 i/s		
cgi:	128455.5 i/s	- 2.23x	slower
gsub:	13743.6 i/s	- 20.84x	slower
rust_map:	7878.5 i/s	- 36.36x	slower
each_char:	4540.9 i/s	- 63.08x	slower

JRuby 9.1:

cgi:	122654.7 i/s		
gsub:	18363.9 i/s	- 6.68x	slower
each_char:	5596.8 i/s	- 21.92x	slower

JRuby 9.1 + 30 s warm-up + invoke dynamic:

cgi:	121956.6 i/s		
gsub:	17826.9 i/s	- 6.84x	slower
each_char:	6063.1 i/s	- 20.11x	slower

MEANWHILE IN CRYSTAL...

```
class String
```

```
  ENTITIES = { '<' => "&lt;", '>' => "&gt;", '&' => "&amp;",  
               '\'' => "&apos;", '"' => "&quot;" }
```

```
  def escape_gsub
```

```
    gsub(ENTITIES)
```

```
  end
```

```
  def escape_html
```

```
    HTML.escape(self)
```

```
  end
```

```
  def escape_io
```

```
    io = IO::Memory.new
```

```
    HTML.escape(self, io)
```

```
    io.to_s
```

```
  end
```

Crystal 0.23 (2017):

io	67.56k	(14.8μs)	(± 1.38%)	fastest
gsub	60.62k	(16.5μs)	(± 2.13%)	1.11× slower
html	57.82k	(17.29μs)	(± 1.58%)	1.17× slower

Crystal 0.24 (2018):

io	35.84k	(27.9μs)	(± 2.84%)	1.18× slower
gsub	42.28k	(23.65μs)	(± 1.05%)	1.00× slower
html	42.36k	(23.61μs)	(± 1.34%)	fastest

Crystal 0.26 (2018):

io	41.23k	(24.26μs)	(± 2.09%)	10816 B/op	1.15× slower
gsub	47.54k	(21.03μs)	(± 2.33%)	9824 B/op	fastest
html	46.49k	(21.51μs)	(± 1.64%)	9744 B/op	1.02× slower

Ruby 2.5:

rust_push:	286450.1	i/s	
cgi:	128455.5	i/s	- 2.23x slower
gsub:	13743.6	i/s	- 20.84x slower
rust_map:	7878.5	i/s	- 36.36x slower
each_char:	4540.9	i/s	- 63.08x slower

FOR SUB-μS TIMES: OPTIMISING STRING PROCESSING IN RUST



andrew dongs

@mustlovedongs

[explaining the theory of relativity to physics students] no no no, einstein is the name of its creator, you mean einstein's monster

BOUNDARY CROSSING

NOT PASSING ANYTHING

BOOLEANS

NUMBERS ($> 2^{62}$? BigIntegers?)

STRINGS (LONGER THAN 23 BYTES?)

ARRAYS (...OF WHAT? CONTAINING self?)

HASHES (...OF WHAT \rightarrow WHAT? WITH self KEY?)

STRUCTS | OBJECTS

RUBY (FRANKEN)GEMS

json

JSON::Pure (RUBY) vs JSON::Ext (C | JAVA)

levenshtein

C WITH TRANSPARENT FALLBACK TO RUBY

two Strings | two Arrays | two [#hash + #eq?]#each

fast_blank

FASTER String#blank?



Settle In, Babies

@alexandraerin

Actually "gif" is the name of the creator. The file format is "gif's file format".

WORLD DOMINATION PLANS

1. BENCHMARK 2. ALGORITHM 3. IMPLEMENTATIONS?

READABILITY, PERFORMANCE, DYNLANG CALL OVERHEAD

READ THE SOURCE! (ESP. FOR SELF-HOSTED LANGUAGES)

NEVER CREATE RUBY STRINGS LONGER THAN 23 CHARACTERS

WRITING RUBY GEMS WITH RUST AND HELIX

HOW TO WRITE RUBY EXTENSIONS WITH CRYSTAL

RUBY AND GO SITTING IN A TREE



Andrew Heister

@andrew_heister

Actually Frankenstein was the name of the scientist.

I, the person correcting you on this trivial point, am the monster.



@chastell

rebased.com

THANKS!

PIOTR SZOTKOWSKI

talks.chastell.net